

# PostgreSQL Performance and Query Tuning Interview Questions & Answers

## Query Analysis & Execution Plans

### Q1: How do you analyze query performance in PostgreSQL?

**Answer:** Use `EXPLAIN` and `EXPLAIN ANALYZE` to understand query execution plans:

```
sql
-- Shows the planned execution without running the query
EXPLAIN * FROM users WHERE age > 25;

-- Actually executes and shows real performance metrics
EXPLAIN ANALYZE SELECT * FROM users WHERE age > 25;

-- More detailed output with buffers, timing, and costs
EXPLAIN (ANALYZE, BUFFERS, TIMING) SELECT * FROM users WHERE age > 25;
```

Key metrics to examine:

- **Cost:** Estimated expense (startup cost..total cost)
- **Rows:** Estimated vs actual rows returned
- **Width:** Average row size in bytes
- **Actual Time:** Real execution time per node
- **Buffers:** Shared/local/temp buffers hit/read/written

### Q2: What are the different types of scans in PostgreSQL and when are they used?

**Answer:**

1. **Sequential Scan (Seq Scan):** Reads entire table row by row
  - Used when no suitable index exists or small tables
  - Cost:  $O(n)$  where  $n$  is table size
2. **Index Scan:** Uses index to locate specific rows

- Efficient for selective queries
- Cost:  $O(\log n)$  for lookup + cost of fetching rows

### 3. Index Only Scan:

- Data retrieved entirely from index
- Possible when all required columns are in the index
  - Fastest for covered queries

### 4. Bitmap Index Scan:

- Creates bitmap of matching rows
- Used for multiple index conditions or low selectivity
  - Combines multiple indexes efficiently

```
sql
```

```
-- Example forcing different scan types
SET enable_seqscan = off; -- Force index usage
SET enable_indexscan = off; -- Force bitmap scan
```

## Q3: How do you interpret and optimize join operations?

**Answer:** PostgreSQL uses three main join algorithms:

### 1. Nested Loop Join:

- Best for small datasets or when one side is very selective
- Cost:  $O(n*m)$
- Often used with indexes

### 2. Hash Join:

- Good for medium to large datasets
- Builds hash table from smaller relation
- Cost:  $O(n+m)$

### 3. Merge Join:

- Efficient for pre-sorted data
- Requires both inputs to be sorted
- Cost:  $O(n \log n + m \log m)$  if sorting needed

Optimization strategies:

sql

-- Ensure join columns have indexes

```
CREATE INDEX idx_orders_customer_id ON orders(customer_id);
```

```
CREATE INDEX idx_customers_id ON customers(id);
```

-- Use appropriate data types (avoid implicit conversions)

-- Correct

```
WHERE customer_id = 123
```

-- Avoid

```
WHERE customer_id = '123' -- if customer_id is integer
```

## Indexing Strategies

**Q4: What are the different types of indexes in PostgreSQL and when should you use each?**

**Answer:**

### 1. B-Tree Index (Default):

sql

```
CREATE INDEX idx_name ON table(column);
```

- Best for equality and range queries
- Supports sorting operations
- Use for: =, <, <=, >, >=, BETWEEN, IN, IS NULL

### 2. Hash Index:

sql

```
CREATE INDEX idx_hash ON table USING HASH(column);
```

- Only equality operations (=)
- Slightly faster than B-tree for equality
- Not WAL-logged (use with caution)

### 3. GIN (Generalized Inverted Index):

sql

```
CREATE INDEX idx_gin ON table USING GIN(column);
```

- For composite values (arrays, JSONB, full-text search)
- Excellent for "contains" operations
- Use for: @>, <@, &&, ?

#### 4. GiST (Generalized Search Tree):

sql

```
CREATE INDEX idx_gist ON table USING GIST(column);
```

- For geometric data, full-text search
- Supports custom data types
- Less precise than GIN but smaller

#### 5. BRIN (Block Range Index):

sql

```
CREATE INDEX idx_brin ON table USING BRIN(column);
```

- For very large tables with natural ordering
- Minimal storage overhead
- Best for time-series or sequential data

#### 6. SP-GiST (Space-Partitioned GiST):

sql

```
CREATE INDEX idx_spgist ON table USING SPGIST(column);
```

- For clustered data with uneven distribution
- Good for text, IP addresses, geographic data

### Q5: How do you create and optimize partial indexes?

**Answer:** Partial indexes include only rows meeting specified conditions:

sql

```
-- Index only active users
CREATE INDEX idx_active_users ON users(email)
WHERE status = 'active';

-- Index only recent orders
CREATE INDEX idx_recent_orders ON orders(created_at)
WHERE created_at >= '2024-01-01';

-- Index only non-null values
CREATE INDEX idx_phone ON users(phone_number)
WHERE phone_number IS NOT NULL;
```

Benefits:

- Smaller index size
- Faster maintenance
- Reduced storage requirements
- Better performance for specific query patterns

Best practices:

- Use when querying a subset of data frequently
- Ensure WHERE clause in queries matches index condition
- Monitor query plans to verify index usage

## Q6: Explain composite indexes and column ordering strategy.

**Answer:** Composite indexes span multiple columns. Column order significantly affects performance:

```
sql
-- Create composite index
CREATE INDEX idx_user_search ON users(status, created_at, email);
```

### Column Ordering Rules:

1. **Most Selective First:** Place columns with highest cardinality first

2. **Equality Before Range:** Put equality conditions before range conditions
3. **Query Pattern Matching:** Order based on most common query patterns

```
sql
```

```
-- Good: equality first, then range
CREATE INDEX idx_orders ON orders(status, customer_id, created_at);

-- Query that uses this index efficiently
SELECT * FROM orders
WHERE status = 'pending'
AND customer_id = 123
AND created_at >= '2024-01-01';
```

### Index Usage Rules:

- PostgreSQL can use leftmost columns of composite index
- Cannot skip columns in the middle efficiently
- Range conditions stop further column usage

## Performance Monitoring

### Q7: What are the key performance metrics to monitor in PostgreSQL?

**Answer:**

#### System-Level Metrics:

```
sql
```

```
-- Database connections
SELECT count(*) as active_connections
FROM pg_stat_activity
WHERE state = 'active';

-- Cache hit ratio (should be >99%)
SELECT
    sum(heap_blks_hit) / (sum(heap_blks_hit) + sum(heap_blks_read)) as cache_hit_ratio
FROM pg_statio_user_tables;

-- Index usage ratio
SELECT
    schemaname,
    tablename,
    round(100.0 * idx_scan / (seq_scan + idx_scan), 2) as index_usage_ratio
FROM pg_stat_user_tables
WHERE seq_scan + idx_scan > 0;
```

## Query-Level Metrics:

```
sql
```

-- Top slow queries (requires pg\_stat\_statements)

**SELECT**

query,  
calls,  
total\_time,  
mean\_time,

**ROWS**

**FROM** pg\_stat\_statements

**ORDER BY** total\_time **DESC**

**LIMIT** 10;

-- Lock monitoring

**SELECT**

blocked\_locks.pid **AS** blocked\_pid,  
blocked\_activity.username **AS** blocked\_user,  
blocking\_locks.pid **AS** blocking\_pid,  
blocking\_activity.username **AS** blocking\_user,  
blocked\_activity.query **AS** blocked\_statement

**FROM** pg\_catalog.pg\_locks blocked\_locks

**JOIN** pg\_catalog.pg\_stat\_activity blocked\_activity **ON** blocked\_activity.pid = blocked\_locks.pid

**JOIN** pg\_catalog.pg\_locks blocking\_locks **ON** blocking\_locks.locktype = blocked\_locks.locktype

**JOIN** pg\_catalog.pg\_stat\_activity blocking\_activity **ON** blocking\_activity.pid = blocking\_locks.pid

**WHERE NOT** blocked\_locks.granted;

## Q8: How do you identify and resolve common performance bottlenecks?

**Answer:**

### 1. I/O Bottlenecks:

sql

```
-- Check table and index sizes
SELECT
    schemaname,
    tablename,
    pg_size.pretty(pg_total_relation_size(schemaname||'.'||tablename)) as total_size,
    pg_size.pretty(pg_relation_size(schemaname||'.'||tablename)) as table_size
FROM pg_tables
ORDER BY pg_total_relation_size(schemaname||'.'||tablename) DESC;
```

-- Solutions:

- - Add appropriate indexes
- - Partition large tables
- - Archive old data
- - Increase shared\_buffers

## 2. CPU Bottlenecks:

sql

```
-- Identify expensive queries
SELECT
    query,
    calls,
    total_time / calls as avg_time,
    rows / calls as avg_rows
FROM pg_stat_statements
WHERE calls > 100
ORDER BY total_time / calls DESC;
```

-- Solutions:

- - Optimize query logic
- - Add indexes
- - Rewrite complex queries
- - Use materialized views

## 3. Memory Issues:

sql

```
-- Check work_mem usage
SHOW work_mem;

-- Monitor temp file usage
SELECT
    schemaname,
    tablename,
    n_tup_ins,
    n_tup_upd,
    n_tup_del
FROM pg_stat_user_tables;

-- Solutions:
-- - Increase work_mem for complex queries
-- - Optimize sort operations
-- - Use streaming replication for read replicas
```

## Configuration Tuning

**Q9: What are the most important PostgreSQL configuration parameters for performance?**

**Answer:**

**Memory Settings:**

```
bash
```

```
# postgresql.conf

# Shared memory for caching (25% of RAM)
shared_buffers = '4GB'

# Memory for sorting and hash operations
work_mem = '64MB'

# Memory for maintenance operations
maintenance_work_mem = '1GB'

# Memory for WAL buffers
wal_buffers = '64MB'

# Effective cache size (50-75% of RAM)
effective_cache_size = '12GB'
```

## I/O and Checkpoints:

```
bash

# WAL settings for performance
wal_level = replica
checkpoint_timeout = '15min'
checkpoint_completion_target = 0.9
max_wal_size = '4GB'
min_wal_size = '1GB'

# Background writer settings
bgwriter_delay = '200ms'
bgwriter_lru_maxpages = 100
bgwriter_lru_multiplier = 2.0
```

## Query Planner:

```
bash
```

```
# Cost parameters (tune based on hardware)
seq_page_cost = 1.0
random_page_cost = 2.0 # Lower for SSDs (1.1-1.5)
cpu_tuple_cost = 0.01
cpu_index_tuple_cost = 0.005
cpu_operator_cost = 0.0025

# Statistics target for better estimates
default_statistics_target = 100
```

## Q10: How do you optimize PostgreSQL for different workload types?

**Answer:**

### OLTP Workloads (High Concurrency, Short Transactions):

```
bash

# Optimize for many concurrent connections
max_connections = 200
shared_buffers = '2GB'      # Lower ratio (15-25% of RAM)
work_mem = '16MB'          # Smaller per-connection
effective_cache_size = '12GB'

# Fast commit settings
synchronous_commit = on
wal_buffers = '64MB'
checkpoint_timeout = '5min'

# Connection pooling recommended (PgBouncer)
```

### OLAP Workloads (Analytics, Large Queries):

```
bash
```

```
# Optimize for complex queries
shared_buffers = '8GB'      # Higher ratio (25-40% of RAM)
work_mem = '256MB'          # Larger for sorting/hashing
maintenance_work_mem = '2GB'
effective_cache_size = '24GB'

# Allow longer operations
statement_timeout = '30min'
lock_timeout = '10min'

# Parallel query settings
max_parallel_workers = 8
max_parallel_workers_per_gather = 4
```

## Mixed Workloads:

```
bash

# Balanced configuration
shared_buffers = '4GB'
work_mem = '64MB'
effective_cache_size = '16GB'

# Use connection pooling
# Separate read replicas for analytics
# Partition large tables by time/range
```

## Advanced Optimization Techniques

### Q11: Explain table partitioning strategies and their performance benefits.

**Answer:**

#### Range Partitioning (Most Common):

```
sql
```

-- Parent table

```
CREATE TABLE sales (
    id SERIAL PRIMARY KEY,
    sale_date DATE NOT NULL,
    amount DECIMAL,
    customer_id INTEGER
) PARTITION BY RANGE (sale_date);
```

-- Partitions

```
CREATE TABLE sales_2023 PARTITION OF sales
FOR VALUES FROM ('2023-01-01') TO ('2024-01-01');

CREATE TABLE sales_2024 PARTITION OF sales
FOR VALUES FROM ('2024-01-01') TO ('2025-01-01');
```

-- Indexes on partitions

```
CREATE INDEX idx_sales_2024_customer ON sales_2024(customer_id);
```

## Hash Partitioning:

sql

-- For distributing data evenly

```
CREATE TABLE users (
    id SERIAL PRIMARY KEY,
    email VARCHAR(255),
    created_at TIMESTAMP
) PARTITION BY HASH (id);
```

-- Create 4 hash partitions

```
CREATE TABLE users_0 PARTITION OF users FOR VALUES WITH (MODULUS 4, REMAINDER 0);
CREATE TABLE users_1 PARTITION OF users FOR VALUES WITH (MODULUS 4, REMAINDER 1);
CREATE TABLE users_2 PARTITION OF users FOR VALUES WITH (MODULUS 4, REMAINDER 2);
CREATE TABLE users_3 PARTITION OF users FOR VALUES WITH (MODULUS 4, REMAINDER 3);
```

## Benefits:

- Partition pruning: Only scan relevant partitions
- Parallel processing: Operations can run on multiple partitions

- Maintenance efficiency: Operations on individual partitions
- Improved query performance for time-based queries

## Q12: What are materialized views and when should you use them?

### Answer:

Materialized views store query results physically and can be refreshed:

sql

-- Create materialized view for expensive aggregation

```
CREATE MATERIALIZED VIEW monthly_sales_summary AS
```

```
SELECT
```

```
DATE_TRUNC('month', sale_date) as month,
```

```
COUNT(*) as total_sales,
```

```
SUM(amount) as total_revenue,
```

```
AVG(amount) as avg_sale_amount
```

```
FROM sales
```

```
GROUP BY DATE_TRUNC('month', sale_date);
```

-- Create index on materialized view

```
CREATE INDEX idx_monthly_sales_month ON monthly_sales_summary(month);
```

-- Refresh strategies

```
REFRESH MATERIALIZED VIEW monthly_sales_summary; -- Blocking
```

```
REFRESH MATERIALIZED VIEW CONCURRENTLY monthly_sales_summary; -- Non-blocking (needs unique
```

### Use Cases:

- Complex aggregations that don't change frequently
- Joining multiple large tables
- Dashboard and reporting queries
- Data warehouse scenarios

### Best Practices:

- Create unique indexes for concurrent refresh
- Schedule refreshes during low-traffic periods

- Monitor refresh times and storage overhead
- Consider incremental refresh strategies

## Troubleshooting Performance Issues

### Q13: How do you diagnose and fix slow queries?

Answer:

#### Step 1: Identify Slow Queries

```
sql
-- Enable query logging in postgresql.conf
log_min_duration_statement = 1000 -- Log queries > 1 second

-- Or use pg_stat_statements
SELECT
    query,
    calls,
    total_time,
    mean_time,
    stddev_time,
    rows
FROM pg_stat_statements
ORDER BY mean_time DESC
LIMIT 10;
```

#### Step 2: Analyze Execution Plan

```
sql
-- Get detailed execution plan
EXPLAIN (ANALYZE, BUFFERS, FORMAT JSON)
SELECT u.name, COUNT(o.id) as order_count
FROM users u
LEFT JOIN orders o ON u.id = o.user_id
WHERE u.created_at >= '2024-01-01'
GROUP BY u.id, u.name;
```

## Step 3: Common Fixes

### 1. Missing Indexes:

sql

-- Add index for WHERE clause

```
CREATE INDEX idx_users_created_at ON users(created_at);
```

-- Add index for JOIN

```
CREATE INDEX idx_orders_user_id ON orders(user_id);
```

### 2. Inefficient JOINS:

sql

-- Bad: N+1 query pattern

-- Fix with proper JOIN and batch processing

-- Bad: SELECT \* with large result sets

-- Fix: SELECT only needed columns

```
SELECT u.name, COUNT(o.id) -- Instead of SELECT *
```

### 3. Suboptimal WHERE Clauses:

sql

-- Bad: Function calls prevent index usage

```
WHERE UPPER(name) = 'JOHN'
```

-- Good: Use functional index or case-insensitive comparison

```
CREATE INDEX idx_name_upper ON users(UPPER(name));
```

-- OR

```
WHERE name ILIKE 'john'
```

## Q14: How do you handle deadlock situations?

**Answer:**

**Understanding Deadlocks:** Deadlocks occur when two or more transactions wait for each other to release locks.

## Detection and Monitoring:

```
sql
-- Check current locks
SELECT
    t.schemaname,
    t.tablename,
    l.locktype,
    l.mode,
    l.granted,
    a.query,
    a.query_start,
    age(now(), a.query_start) AS duration
FROM pg_locks l
JOIN pg_stat_activity a ON l.pid = a.pid
JOIN pg_tables t ON l.relation = t.schemaname || '.' || t.tablename::regclass
ORDER BY a.query_start;

-- Deadlock detection settings
deadlock_timeout = '1s' -- How long to wait before checking for deadlock
```

## Prevention Strategies:

### 1. Consistent Lock Ordering:

```
sql
-- Always acquire locks in the same order across transactions
BEGIN;
LOCK TABLE users IN SHARE MODE;
LOCK TABLE orders IN SHARE MODE;
-- ... perform operations
COMMIT;
```

### 2. Shorter Transactions:

```
sql
```

```
-- Keep transactions brief and avoid long-running operations
BEGIN;
-- Quick operations only
UPDATE users SET last_login = NOW() WHERE id = 123;
COMMIT;
```

### 3. Appropriate Isolation Levels:

```
sql
-- Use appropriate isolation level
SET TRANSACTION ISOLATION LEVEL READ COMMITTED; -- Default, usually sufficient
-- Avoid SERIALIZABLE unless necessary
```

#### Resolution:

- PostgreSQL automatically detects and resolves deadlocks
- One transaction is aborted (deadlock victim)
- Application should retry the aborted transaction
- Monitor deadlock frequency and patterns

## Best Practices Summary

### Q15: What are the top 10 PostgreSQL performance best practices?

#### Answer:

#### 1. Index Strategy:

- Create indexes on frequently queried columns
- Use composite indexes for multi-column queries
- Monitor and remove unused indexes
- Use partial indexes for subset queries

#### 2. Query Optimization:

- Use EXPLAIN ANALYZE to understand query plans
- Avoid SELECT \* in application code
- Use appropriate WHERE clauses with indexed columns

- Optimize JOINs with proper indexing

### 3. Configuration Tuning:

- Set shared\_buffers to 25% of RAM
- Configure work\_mem based on concurrent users
- Set effective\_cache\_size to 50-75% of RAM
- Tune checkpoint settings for workload

### 4. Table Design:

- Use appropriate data types (smaller is better)
- Normalize appropriately (avoid over-normalization)
- Consider partitioning for large tables
- Use constraints to help query planner

### 5. Maintenance:

- Regular VACUUM and ANALYZE operations
- Monitor table bloat and reindex when needed
- Update table statistics regularly
- Archive old data to keep tables manageable

### 6. Connection Management:

- Use connection pooling (PgBouncer, pgpool)
- Set appropriate max\_connections
- Monitor connection usage
- Close connections properly in applications

### 7. Hardware Optimization:

- Use SSDs for better random I/O
- Separate WAL files on different disks
- Ensure adequate RAM for caching
- Configure OS-level parameters

### 8. Monitoring:

- Enable pg\_stat\_statements extension
- Monitor slow query log

- Track key performance metrics
- Set up alerting for performance degradation

## 9. **Backup and Recovery:**

- Use streaming replication for high availability
- Configure WAL archiving properly
- Test recovery procedures regularly
- Monitor replication lag

## 10. **Application-Level Optimization:**

- Implement proper error handling and retries
- Use batch operations for bulk data changes
- Cache frequently accessed data at application level
- Design database schema with query patterns in mind